EPRO/IAS

Bachelor Studiengang Medientechnik
Fachhochschule St.Pölten

Patrik Lechner

Wien, January 12, 2017

1. Introduction

Figure 1.1.: Weird effects of digital signals in video

This section tries to make sure we are all on the same page. It goes through some mathematics notation and principles of digital audio.

If you rather want a video format revision of digital audio (I mean, it's 2017, of course you want), I can recommend D/A and A/D | Digital Show and Tell (Monty Montgomery @ xiph.org) [1]. Some of the points in the video go beyond what we are doing here and vice versa, so don't purely rely on the video though.

Ideally, if you throughly understood the introduction section, you should be able to go through the remaining chapters with tempo.

1.1. About this document

Please report any mistakes, errors etc to ptrk.lechner@gmail.com.

Some information in this document is relevant for understanding its contents but not relevant for the exam. For example in chapter 2 we will find a really complicated result of an equation. The point there is just that it's complicated. And it is shown how complicated, but this is nothing to be learned by heart. Importance is tried to be made clear using the following formats:

> Text, figures, and equations with a gray background like this are background information that is not to be learned by heart.

Video Analogies
This document tries to explain digital signals. It does this by use of audio signals mainly. Sometimes video analogies are given. These are also not relevant for the exam.

Very important things are framed.

pd objects will always be in courier new and square brackets. Like this: `[metro]`

1.2. About plotting signals

We will need to plot a lot of signals in order to understand them better. Most of the time, such a plot will look like figure 1.2.

---

[1] https://www.youtube.com/watch?v=cIQ9IXSUzuM

Figure 1.2.: Sine wave, 1Hz, sampled at 30Hz sample rate

This plot looks nice but it has a problem. The sine wave is sampled at a sampling rate of 30 Hz, but we see a continuous line. This "connection of the dots" is created by plotting. It is kind of similar to what our *digital-to-analog-converter* does. It somehow[2] interpolates the values we have.

This can be misleading, so we should actually plot something more like figure 1.3. Often we can see plots that look like figure 1.4 as well when a signal is analyzed.

[2]linear interpolation in case of the plot

Figure 1.3.: Sine wave, 1Hz, sample rate 30Hz, scatter-plot

Figure 1.4.: Sine wave, 1Hz, sample rate 30Hz, stem-plot

So why don't we always do a stem- or scatter-plot? Simply because it gets too crowded with our usual sampling rates in audio. It just works with very low sampling rates or very short signals.

Figure 1.5.: Sine wave, 1Hz, sample rate 44100Hz, stem-plot

But we should never forget that we don't actually have the values in between the dots. Digital signals are not defined between their sampled points.

1.3. What is aliasing?

Aliasing in audio means problems caused by signals that exceed the nyquist rate. The nyquist rate, let's call it f_n for now, is defined by the half of the sample rate (f_s). So,

$$f_n = \frac{f_s}{2} \tag{1.1}$$

A digital system can only describe signals up to his nyquist rate. If we try to make signals higher than this frequency, we will fail and encounter strange effects.

Visually speaking, frequencies higher than nyquist fold back. So, let's assume we have a sampling rate of 100Hz. Nyquist would be at 50Hz. If we try to synthesize a sine wave with 51Hz, what we will get is a 49Hz one. If we try to make a 52Hz one, we will get 48Hz. So you see, it simply folds back.

Video analogies

Aliasing in graphics usually means *spacial* aliasing, so aliasing in the space domain. This is what we see in figure 1.6. But there is also time domain aliasing in film. It is actually more natural to think about the sampling rate in audio as the same as the frame rate in video. For some really weird effects that arise in video due to time domain aliasing see airplane[3], Water experiment[4] or Guitar strings[5].

Figure 1.6.: Spacial aliasing in graphics. The "frequency" of the intended pixels is to high for the actual pixels.

The particularly strange effects in the airplane example above are caused by the rolling shutter of a CMOS sensor. Since it is not sampling the incoming light uniformly (at the same time) the image is distorted.

In fugure 1.7 you can see a visualization of aliasing in the time domain.

[2]https://www.youtube.com/watch?v=LVwmtwZLG88
[3]https://www.youtube.com/watch?v=GBtHeR-hY9Y
[4]https://www.youtube.com/watch?v=jcOKTTnOIV8

Figure 1.7.: Aliasing visualized in the time domain. A sampling rate of 4 Hz is used. Therefore frequencies will fold above 2 Hz, which is the nyquist frequency. The input cosine has a frequency of 3 Hz, labeled "Orginal Signal". What we would get is the sampled points. If these are digital-to-analog converted, we would get what is labeled "Aliased signal " in this plot, so a 1 Hz cosine.

1.4. Scaling and Mapping Signals

It is an important skill to be able to scale signals from one range to another. We need it a lot and we will be able to think about signals more easily if we mastered this task. It's actually quite simple, we just have to imagine the signals visually.

So what exactly do we have to do here? We are confronted with the following problem: Given some signal, say, a sine wave with its maximum at the value 1 and its minimum at the value -1. How to bring it to a different range, say, 0-10?

It helps me a lot to solve this problem in two parts: first get the input in the range 0-1, then from there go to the desired range. What can we do to the signal? Let's take a sine wave:

Well we can add and subtract to move the wave vertically, so let's add 1 to move it up:

So we can move signals around by adding constant values. We can scale them by

Figure 1.8.: a sine wave

Figure 1.9.: the same sine wave, 1 added to a each sample, therefore shifted upwards.

multiplication. So if we take our sine that now ranges from 0 to 2 and multiply it by 0.5, we get whats in figure 1.10.

Using what we got now in figure 1.10, we can just multiply by 10, easy! Beware that there are always multiple solutions to this kind of a problem. Try to find another one for the problem above!

Figure 1.10.: Sine with a range of 0-1. Obtained by taking a sine wave, adding 1 and dividing by two afterwards.

> **Question 1** *Let's take a sine wave that has it's minimum at 2 and it's maximum at 5. What do we have to do to get it into a -1 to 1 range?*

1.5. What's DC-Offset?

What we did above by adding a constant value to a signal can be called adding DC-offset ("Gleichspannungsversatz"), DC-Bias or a DC component. These are different words for the same thing.

DC-offset can also be encountered in signals we recorded (caused by old or broken equipment mainly). But we have seen that we can also generate DC-offset.

To state it again clearly:

> DC-Offset is a constant value over time, or a constant offset from the zero value in Y. So for example in figure 1.10 we see a DC-offset of 0.5 since subtracting this value would center the wave around 0.

If we try to think how this kind of signal looks in the frequency domain we find that it is energy at 0Hz. In figure 1.11 you can see a couple of cosine waves plotted. This

should hep imagine that a constant signal, a signal that does not move a all, can be described by a cosine with 0 Hz and a certain amplitude.


Figure 1.11.: Cosines with different very low frequencies approaching 0Hz, so DC offset.

1.6. What's an Impulse?

work in progress.

1.7. How to describe audio mathematically

If we want to talk about signals, or if we want to analyze them, it is often useful to look at the problem mathematically. First, let's introduce some conventions. They might look unfamiliar or complicated. But in fact it is not very complicated and knowing these conventions makes it easy to communicate (e.g. reading scientific papers about our topics or explaining something to another person).

1.7.1. signals

Usually we describe a *digital* signal by a name, say x, (but you can call it however you want). If we want to talk about the individual samples, or values of the (mono) signal, we can use a subscript or parenthesis. So if the fith sample of x is 1, we could write:

$$x_5 = 1 \tag{1.2}$$

or

$$x(5) = 1 \tag{1.3}$$

Oftentimes we like to talk about a signal more generally and we use n as a place holder for this index, so we might write x_n, meaning the nth sample of x.

sine waves

We use sine waves a lot. The syntax can get a bit overwhelming at first, so let's quickly explain what's going on in a standard sine wave oscillator.

$$x(t) = A \cdot sin(2\pi ft + \phi) \tag{1.4}$$

Where f is the frequency in Hertz and t is the time in seconds. ϕ Is a (possibly constant) phase offset and A can be used to scale the whole thing. Since cosine and sine have their peaks at -1 and 1, A will be the amplitude. If A is set to 0.5, the resulting signal will have it's peaks at -0.5 and 0.5.
The upper equation is rather complete. Most of the time, we will just ignore the phase and the amplitude and just write:

$$x(t) = sin(2\pi ft) \tag{1.5}$$

Sometimes, we will simplify even more and just write $sin(a)$ or $sin(b)$ or similar.

In literature, we sometimes encounter $x(t) = sin(t \cdot \omega)$. ω stands for a "frequency"(actually for the rotational speed) in *radians per second*. 2π radians per second are 1 Hz and we can therefore convert radians per second to frequency, v, with:

$$v = \omega/2\pi \tag{1.6}$$

This notation is only covered since it is very common, but it will not be used here, since it is considered more intuitive to work with frequency in Hz.

What is actually the significance of using either *sin* or *cos*?
In most cases for us, this does not matter at all. The difference between sine and cosine is just in phase. Since we are most of the time describing a sine (or cosine) wave oscillator as a function of time, the only difference between the two would be their phase if given some moment in time. This is nothing one could hear also.

1.7.2. systems

There are multiple ways in which we can describe linear time invariant (LTI) systems. We will talk more about this in chapter 3. Here, we will just look at how *difference equations* work. These are one possibility to describe a digital LTI system. They are the discrete equivalent of difference equations. It's really not that complicated:

$$y(n) = x(n) \tag{1.7}$$

Would be an equation that describes a system that does nothing. it takes the input sample, $x(n)$ does nothing and defines the output sample $y(n)$ with it. Another really simple system would multiply its input by two:

$$y(n) = x(n) \cdot 2 \tag{1.8}$$

It's getting interesting if we start playing with the index:

$$y(n) = x(n - 1) \tag{1.9}$$

Is a delay by one sample.

$$y(n) = x(n) + x(n - 1) \tag{1.10}$$

Takes it's input sample and the previous input sample, adds them together and spits it out. Can you imagine what this does? We'll get to it in chapter 3...
This is just a preview, but it really gets interesting if we start working with feedback:

$$y(n) = x(n) + y(n - 1) \tag{1.11}$$

1.8. Message Domain/Signal Domain

This is not pure data specific although it might sound like it. In pure data we can have audio signals. Audio signals are processed in buffers. They are just numbers, but these numbers are calculated at a rate of 44100 Hz *all the time* if we choose to have our sample rate at 44100 Hz.
Messages on the other hand are not calculated that often and not all the time. Messages are processed *on demand*. They are event based. This means that, for example, if we hit a note on our midi keyboard or enter a number in a `[Numberbox]` this *event* will pass through its following objects *once*. Have a look at figure 1.12 or at the patcher. Also note that pd helps us in understanding if we deal with message or signal domain:

- pd is indicating the signal domain with thicker patch cords

- signal domain inlets/outlets are back (look closely)

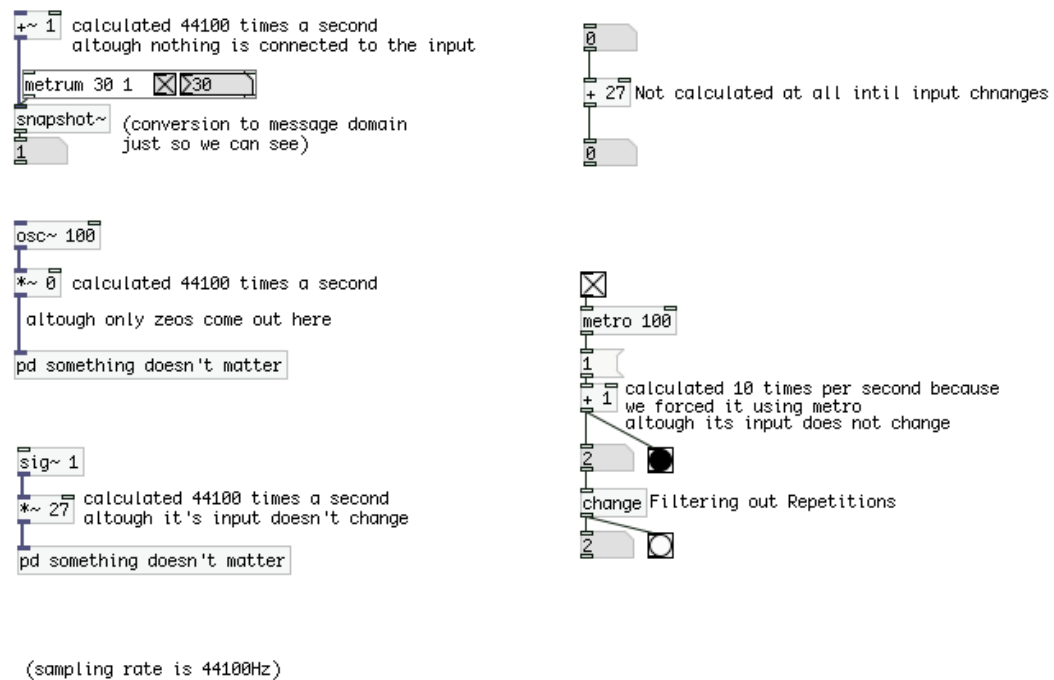- objects that deal with the signal domain have a tilde (~) at the end of their name.


```
+~ 1   calculated 44100 times a second
       altough nothing is connected to the input

metrum 30 1  ☒ ▷30

snapshot~   (conversion to message domain
1           just so we can see)


osc~ 100

*~ 0   calculated 44100 times a second

       altough only zeos come out here

pd something doesn't matter


sig~ 1

*~ 27  calculated 44100 times a second
       altough it's input doesn't change

pd something doesn't matter


(sampling rate is 44100Hz)
```

```
0

+ 27  Not calculated at all intil input chnanges

0


☒
metro 100

1
+ 1   calculated 10 times per second because
      we forced it using metro
      altough its input does not change
2     ●

change  Filtering out Repetitions
2       ○
```

Figure 1.12.: The patcher `messageDomainSignalDomain.pd` should demonstrate the differences between message domain and signal domain in pure data.

How is this not specific to pure data? The key points here are not specific to this programming language. If we, for example adjust the volume of a track in pro tools or ableton live or similar, then the information of our fader movement will also be in some kind of message domain.

One key aspect here is: some kind of conversion between the two is necessary if we want to switch between them. If we want to control the volume of an audio signal from the message domain we have a problem: we will get noisy output since the message domain is not running at sample rate. Look at a naive attempt to controlling the amplitude of a sine wave:

And let's look at what kind of waveform it will produce:

Figure 1.13.: The patcher `amplitudeWrongRight.pd`

Figure 1.14.: This is a cosine wave at 20 Hz and the message domain running at 100 millisecond interval to make the effect more extreme and visible

This is a problem, since the sudden jumps produce clicks, crackling or zipper noise

(reproduce or open the patcher and hear it for yourself!). These clicks are caused by the fact that the message domain just doesn't produce the numbers at the rate of the sampling rate. It's too slow, so to speak.

In pure data (but also in general) this problem can be solved by interpolation. pure data offers the `[line~]` object for this situation. It can be used to create signals at sampling rate and we can tell it to ramp to a specified value in a given amount of time. In figure 1.15 we see it ramping to any value that comes from the `[/ 127]` object within 20 milliseconds[6]. This interpolation reduces the clicks and noise dramatically. Note that the output of `[line~]` is in the signal domain (thick cable).

Figure 1.15.: Controlling the amplitude of an oscillator. The interpolation done with `[line~]` suppresses clicks.

1.9. Key Points

- Make sure you understand and know the mathematical notation. It will be how we talk in the remaining chapters.

- Make sure you understand what aliasing means in audio.

- Make sure you are comfortable with basic mathematical operations on signals, such as adding constants and multiplying with constants

- Make sure you understand the difference between message and signal domain and the problem with using a message domain "signal" controlling an audio stream.

---

[6]20ms is just an example. It is a reasonable time for this kind of interpolation but really, it is just an example. It could just as well be 50 ms.

Contents

Part I.

Semester 3

1. Sampling, waveshaping, and non-linearity

Figure 1.1.: Lecture Contents

1.1. Waveshaping

Wikipedia quote, page "wavesahper":

„The mathematics of non-linear operations on audio signals is difficult, and not well understood." Waveshaping means distortion. It adds overtones, take a look at figure 1.2.

Figure 1.2.: A sine wave has been generated and waveshaping was applied to add overtones.


Figure 1.3.: The same as the spectrogram above, but in the time domain. We can see the input sine wave and the slightly distorted output. It may look like just the amplitude has changed, but the sine's actual *shape* has changed slightly

1.1.1. The simplest case: a linear Transfer function.

See 1.4. A linear transfer function is used as a lookup table for a sinosoidal input.

Figure 1.4.: Linear Transfer function

A transfer function in the sense of a waveshaper (a "transfer function" might also mean frequency response in other contexts) is a simple look-up function. Waveshaping means to use an input wave to *look up* values in a table or function. A linear transfer function, let's call it l, can result in no change, for example, it might return $l(x) = x$. This means, that whatever value we pass in, we get the same value out. Other linear transfer functions might *only* change the amplitude. For example $f(x) = x \cdot \frac{1}{2}$. That doesn't seem very interesting. But it might explain the term "linear". A transfer function is linear if it looks like a line if we plot it. Look at figure 1.5.

Non-liner transfer functions behave differently. They map their input to other values, such as $f(x) = x^2$. And if we plot then, they don't look like a line. You can also look at figure 1.6 in order to understand what's happening. We again see a linear transfer function but also a non linear one.

Figure 1.5.: A couple of linear transfer functions and their corresponding effects demonstrated using a sine wave. From left to right: multiplication by 0.5, so attenuation by about 6dB, inversion, and the "do-nothing"-function.

Figure 1.6.: A waveshaping visualization taken from Farnell (2010)

But let's get back to our square function, since it's simpler and we will find some surprising results when analyzing it. Let's first simply plot it too.

Figure 1.7.: The function $f(x) = x^2$

Video Analogies

What does applying a curve to every pixel of a video or image do? Is this even something people do? Of course! Think of contrast curves, below is an example from photoshop. Beware that video is working with unipolar input values and audio with bipolar inputs.

Clicking the On-image adjustment tool in the image adds control points to the curve. Moving the control points adjusts the tonality of the image.

Figure 1.8.: Contrast curve in photoshop

1.1.2. Simple non-linearity: X^2

So let's analyze what happens if we use this function for waveshaping. Here it is again:

$$f(x) = x^2 \tag{1.1}$$

Let's simply listen to what's happening, building it in pd:

Figure 1.9.: The square function in pure data, using a 100Hz sine as a test signal. What do you hear?

And we can simply plot what happened if we apply the function before also trying to understand analytically:

Figure 1.10.: Applying the square function to an input sine wave.

Weird, the input seems to double in frequency (did you hear that?). Let's try to understand what's happening.

So we calculate what happens if we send a cosine through this function, so let's take:

$$x = cos(\omega) \tag{1.2}$$

wit arbitrary ω. We can just ignore ω here for a while. Usually, there should be some indexing variable in the cosine function if we want to describe an oscillator that moves over time, but let's also skip that.

So applying our square function we of course get:

$$y = cos(\omega)^2 \tag{1.3}$$

This again results in:

$$y = cos(\omega) \cdot cos(\omega) \tag{1.4}$$

So far so trivial. Note that a multiplication of two oscillators is called *Amplitude Modulation* (actually, in this case we encounter "Ring Modulation", but let's ignore that also), and we know things about Amplitude modulation, namely:

> When multiplying two oscillators, we get sum and difference of the two input frequencies. (And the whole output is attenuated by 6dB)

The above statement in equation form:

$$cos(a) \cdot cos(b) = \frac{cos(a+b) + cos(a-b)}{2} \tag{1.5}$$

We could also have looked up this *trigonometric identity*. This means for our experiment with our cosine squared:

$$y = \frac{cos(\omega + \omega) + cos(\omega - \omega)}{2} \tag{1.6}$$

So:

$$y = \frac{cos(2 \cdot \omega) + cos(0)}{2} = \frac{cos(2 \cdot \omega)}{2} + \frac{1}{2} \tag{1.7}$$

We arrive at the same result! But is this true for every input? That would mean we just built a frequency shifter, did we? No. Waveshaping is much more complicated.

This is immediately obvious when we try to do the same with two oscillators:

$$x = cos(\omega_1) + cos(\omega_2) \tag{1.8}$$

then

$$y = (cos(\omega_1) + cos(\omega_2))^2 \tag{1.9}$$

$$y = cos(\omega_1)^2 + cos(\omega_2)^2 + 2 \cdot cos(\omega_1) \cdot cos(\omega_2) \tag{1.10}$$

And finally:

$$y = \frac{cos(2 \cdot \omega_1)}{2} + \frac{1}{2} + \frac{cos(2 \cdot \omega_2)}{2} + \frac{1}{2} + 2 \cdot (\frac{cos(\omega_1 + \omega_2) + cos(\omega_1 - \omega_2)}{2}) \tag{1.11}$$

1.1.3. How can waveshaping be implemented?

Take a look at figure 1.11. What do you think is happening? On the left side, we see waveshaping as we did it above, using a mathematical function, in this case the tangens hyperbolicus, to distort our signal. On the right side, we see a table that contains the authors desperate attempt to draw the same function with the mouse. The results are theoretically equivalent (if the function in the table was correct), but what are the advantages and disadvantages of the two approaches? Also, be sure to understand what the addition of 1 and the multiplication with 50 does on the right side. Hint: the array has 100 points.
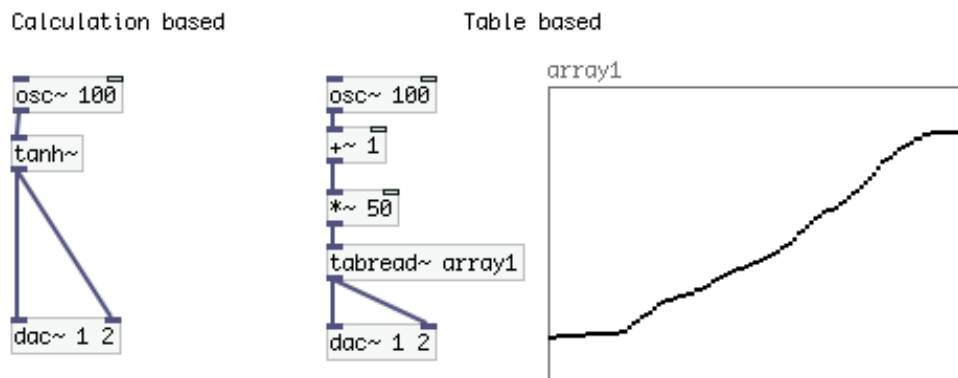
Figure 1.11.: Left: using a mathematical function to calculate the output. Right: using a table to look up the output.

1.1.4. How is Waveshaping related to other techniques?

Sampling

If we take a look at figure 1.19, we see that we play a sound file by accessing a buffer (wavetable) using an index, an oscillator. This is effectively the same setup as we would build for distorting an input sound. Also take a look at figure 1.12, which showing us that waveshaping and wavetable synthesis are identical.

Figure 1.12.: Picture taken from Farnell (2010), showing the identity of waveshaping and wavetable synthesis

Modulation

While we will talk about modulation in a separate chapter, let's loosely define amplitude modulation (AM) as the multiplication of two oscillators and frequency modulation (FM) as varying the frequency of an oscillator using another oscillator. So, as we have also seen above, AM looks like this:

$$y = sin(a) \cdot sin(b) \tag{1.12}$$

and FM looks like this:

$$y = sin(sin(a)) \tag{1.13}$$

in practice, the a and b terms are a bit more complicated, but we will look at this later. That certain cases of AM are identical to waveshaping has been shown above, think

about the square function again. This of course does not mean that waveshaping can do everything AM can do and it does not mean that AM can achieve everything that waveshaping can. This should just show that we can understand the techniques from the perspective of another. What about FM? Well if our lookup function we use for waveshaping is a sine wave, we arrive at the exact same equation as how we defined FM above. Again, practically speaking, the results we get with these two techniques are very different, but we can see the connections.

1.1.5. Why is Waveshaping useful?

The output spectrum is dependent on the input amplitude. This makes it easy to create complex evolving spectra.

1.1.6. What are the problems with waveshaping?

Waveshaping adds overtones. When we build a waveshaper, we have to be aware of aliasing. Take a look at figures 1.13 and 1.14. Sinewaves have been amplified and clipped here.

Figure 1.13.: A sine wave was generated and clipped. Clipping is a form of waveshaping which adds many overtones. Note how high frequencies fold back into the lower parts of the spectrum because they exceed the Nyquist-rate.

In pd we could achieve this like in figure 1.15.
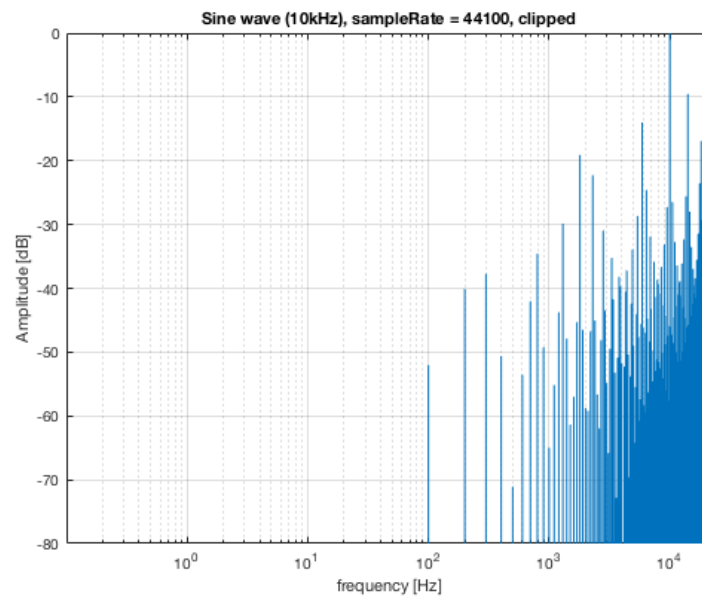
Figure 1.14.: Again, a sine wave, this time with a higher frequency to begin with. Extreme clipping has been applied by boosting the input amplitude. The aliased overtones are all over the place, even below the input frequency.

Figure 1.15.: Clipping an amplified sine wave in pd

What does the output look like? Let's not only look at the spectra but also at the time signal:

Figure 1.16.: An amplified and clipped sine wave in the time domain.

We see that we can arrive at a square-wave like result, but this square-wave is not anti-aliased.

The problem of aliasing in waveshaping is usually treated by over-sampling. This does not solve the problem but lessens it significantly resulting in cleaner, arguably better sound. Oversampling means that, if we work at a sample-rate of 44.1kHz, the input is up-sampled, essentially interpolated, to be at a sample-rate of 88.2kHz. Then the waveshaping is applied, leaving room for high frequencies up to 44.1kHz. Using a lowpass filter, high frequencies over 22.1kHz are then attenuated as much as possible, in order to be able to down-sample again to reach our initial sample-rate of 44.1kHz. To state it more simple: Waveshaping is usually encapsulated in a process that runs at higher sampling rates in order to lessen aliasing.

1.2. Sampler

In pure data (but also in general) we can choose to play a sound from RAM or more or less[1] directly from the hard disk. Both Approaches have their different use-cases.

---

[1]the data played from disk needs to be cached also

Video Analogies

In video, we need to make the same decision essentially. We can choose to buffer as much as we can on the VRAM of the graphics card or rather play from disk. Also different codecs specialize on balancing the load differently. The HAP codec for example tries to only push as much of the load to the hard disk and as little as possible to the CPU (which is decoding the video for playback). In times of SSDs this is a very good (and scalable) method to push a lot of data to playback in realtime.

Work in progress.

Figure 1.17.: simpleSampler

Figure 1.18.: sound in Ram

Loading Audio to an Array (to RAM)

○ Open File

openpanel

t s b

symbol aSoundArray

pack s s

prepend read -resize

soundfiler

124439   # Samples

s numSamps

r numSamps

/ 44100

expr 1/$f1

phasor~

r numSamps

*~

tabread4~ aSoundArray

dac~ 1 2

aSoundArray
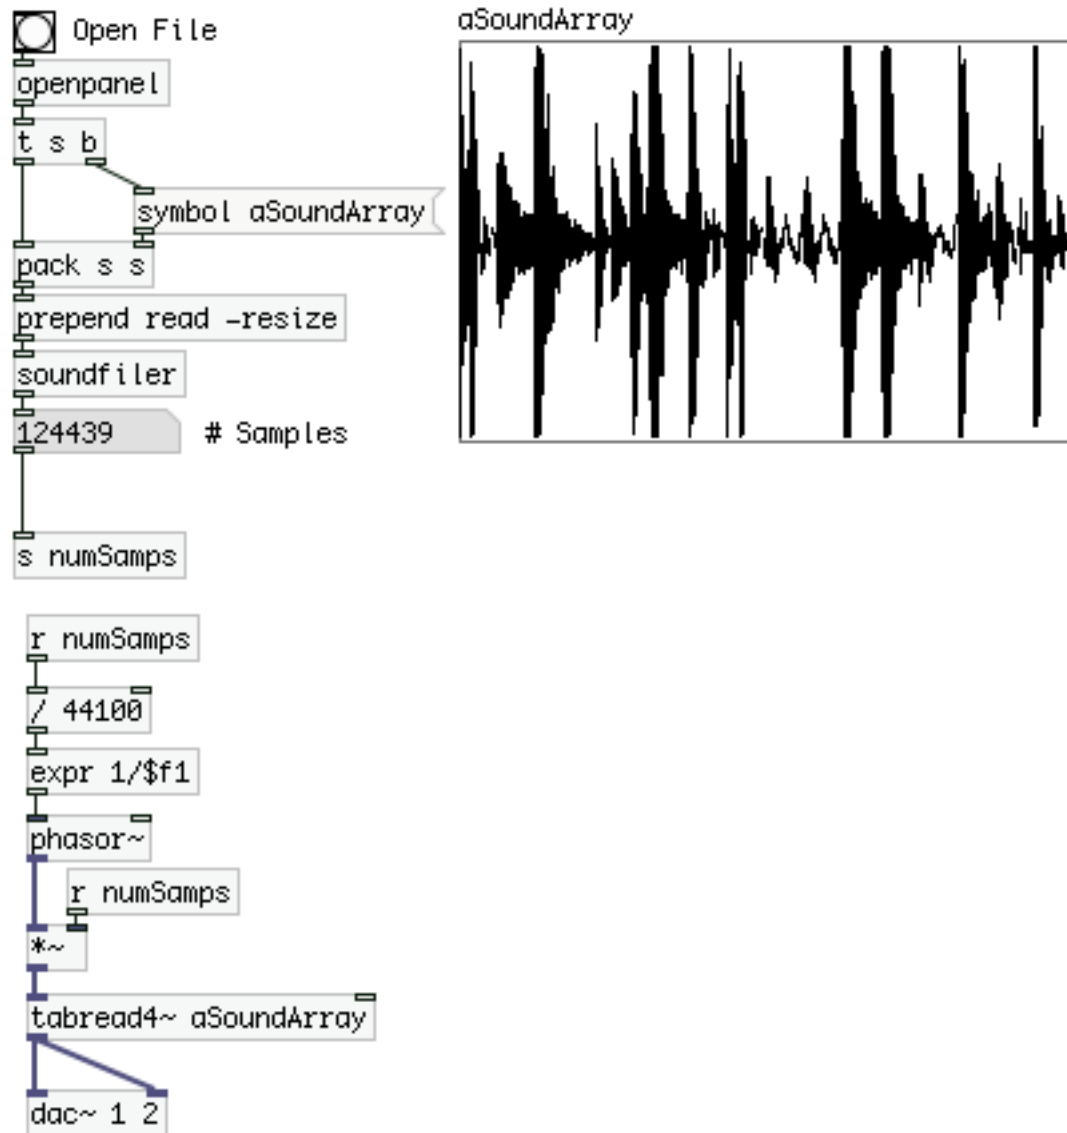
Figure 1.19.: RamFilePlayback

Figure 1.20.: writing Audio to disk

1.2.1. Granular Sampling

In figure 1.21 you see the basic structure of granular sampling. Granular sampling allows us for example to playback a file slower without lowering the pitch. It is one of the methods to make speed and ppitch independent from each other during playback. Other methods also include FFT based approaches.

Figure 1.21.: moreSamplimg.pd, a simplified version of granular sampling

1.3. Key Points

- When do we sampling from RAM when from disk?

- what are the problems with waveshaping?

- Make sure you understand that waveshaping can produce overtones.

- Make sure if you see a (simple) mathematical expression you can decide if it is non-linear or not.

- Make sure you recognize waveshaping in pd if you see it (e.g. figure 1.9)

1.4. Hausübung

1.4.1. Testmodul

baue ein audio Testmodul mit folgener spezifikation:

- Ein audio output

- verschiedene klangquellen wählbar:
    1. White Noise
    2. Sinus (freq. einstellbar)
    3. soundfile (file wählbar)

- GUI

- verfügbar(in eurem pfad, und jederzeit abrufbar als abstraction)

- output pegel sichtbar (level meter)
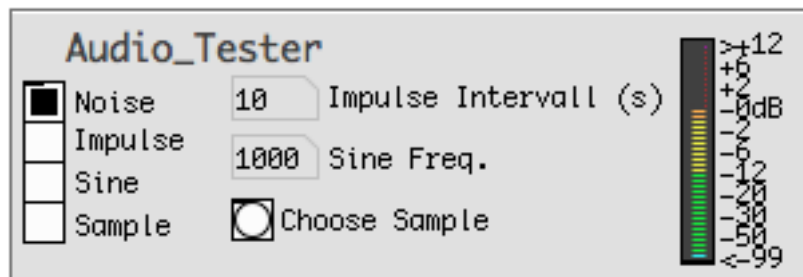
Figure 1.22.: audioTester.pd, zu bauen als Hausübung

2. Modulation

Figure 2.1.: "Surface Modulation" by Richard Sweeney

Figure 2.2.: Lecture Contents

2.1. AM

In this section we will try amplitude modulation. Amplitude modulation is used for radio communication (so you'll need to understand this as a technician, modulation techniques are extremely important and this is maybe the simplest one) but also in sound design. We will try to understand the problem from different perspectives at once:

- Doing some math

- listening to it

- brining it in context to beating waves

- seeing it as convolution in the frequency domain

Amplitude modulation means modulating the amplitude of a signal(surprise!). Modulating means changing over time by another signal. So we have some signal, say, a sine wave, and change its amplitude with another signal, say, another sine wave. Talking this concept and reducing it radically, we end up with figure 2.3.

Figure 2.3.: Simplest form of "Ring modulation".

Of course, we get no sound in figure 2.3 because the frequencies are not initialized, but it shows the general principle. The caption of that figure says "Ring Modulation". Let's quickly get some vocabulary straight:

- "Amplitude modulation" might mean any modulation of amplitude

- "Ring Modulation" means *bi-polar* amplitude modulation.

- In sound design, "Amplitude Modulation" might specifically mean unipolar amplitude modulation.

And some more vocabulary to put what we are doing in a musical context:

- Modulating the amplitude is called "Tremolo" in music [1]

- Modulating the pitch or frequency ("FM") is called "vibrato" in a musical context.[2]

Enough words, let's look at what AM looks like, look at figure 2.4.

> **Question 2** *If we would listen to the signal depicted in the bottom plot of figure 2.4, what do you think we would hear? Try to imagine! If you can't, use pure data to test it! That's why we are using pd.*

> **Answer 2** *We would hear a 30Hz sine wave repeatedly rising and falling in amplitude.*

---

[1]sadly, the fender stratocaster's "Tremolo Arm" is used to control the pitch. Ignore Fender, they got it wrong. You can trust that most guitar players are confused because of this.

[2]Maybe think about it like this: The F in FM is a bit like the v in vibrato. Just to avoid confusion..

Figure 2.4.: Looking at AM in the time domain

Question 3 *Next question, same plot, same signal. So hopefully you found out that we hear a 30 Hz sine with rising and falling amplitude. At what frequency does the amplitude rise and fall? Remember we are modulating with a 1 Hz sine.*

Answer 3 *Two Hertz.*

Maybe you remember from the waveshaping chapter that we actually calculated what frequencies should come out of AM. Also maybe you remember that:

Amplitude Modulation produces sum and difference frequencies of the input frequencies.

$$cos(a) \cdot cos(b) = \frac{cos(a+b) + cos(a-b)}{2} \tag{2.1}$$

But here we still hear the 30 Hz, just getting louder and softer, so what's wrong?
So, let's calculate this. We have two oscillators, $x_1(t) = cos(30t2\pi)$ and $x_2(t) = cos(t2\pi)$. We multiply them, ending up with:

$$y(t) = cos(30t2\pi) \cdot cos(t2\pi) \tag{2.2}$$

Ok, the above formula tells us this means:

$$y(t) = \frac{cos(30t2\pi + t2\pi) + cos(30t2\pi - t2\pi)}{2} \tag{2.3}$$

We can now simplify to:

$$y(t) = \frac{cos((30+1)t2\pi) + cos((30-1)t2\pi)}{2} = \frac{cos(31t2\pi) + cos(29t2\pi)}{2} \qquad (2.4)$$

Hm, so we get out a 31Hz and 29Hz oscillator. What about that rising and falling in amplitude that we hear *and* observe in the plot, surely there must be something wrong! Do you have a solution to this?

Let's use pure data to help us understand. Make two oscillators, one with 29Hz and one with 31Hz, what do you hear?

Figure 2.5.: Adding two oscillators with frequencies very close to each other

If you listen to what is depicted in figure 2.5, you will in fact hear the same as if you do the 30Hz /1 Hz amplitude modulation (which indicates that our formula above is correct). What we hear is a phenomenon called beating. The phase of the two oscillators is canceling each other out at regular intervals because the frequencies are so close to each other. In fact the frequency of the beating f_{beat} is always

$$fbeat = |f_1 - f_2| \qquad (2.5)$$

So the difference of the two frequencies.

Figure 2.6.: caption

Convolution   We will look at convolution in a later chapter. If you are eager to
know what convolution exactly does you can have a look at this[a] . But why mention
convolution if it is not explained at this point? Amplitude modulation and convo-
lution just have a very close relationship, an that is: **Multiplying two signals in
the time domain is equivalent to convolution in the frequency domain
and convolution in the time domain is equivalent to multiplication in
the frequency domain.**

[a]https://youtu.be/_vyke3vF4Nk?t=25m14s

2.2. FM

Frequency modulation can be used to modulate the frequency, meaning to vary the
pitch of a sound, see figure 2.10. But it can also be used to generate overtones and
an overall richer spectrum, see figure 2.11. The two "different versions" only differ
from each other by different parameter[3] values being used. We can think of Frequency
Modulation (FM) as something of the form:

$$y(t) = cos(cos(b)) \tag{2.6}$$

This really is a simplification, but the overall structure of the formula is correct. Look-
ing at figure 2.7 we can see the very basic idea implemented in pure data.

---

[3]We see which parameters can be adjusted in a minute. But if you want to know them now: modulation
frequency, modulation amount and carrier frequency.

The General Idea

Figure 2.7.: The General Idea of FM

Since the frequencies of the oscillators are not set, we won't hear anything when building this patcher. But looking at it may reveal that the idea simply is to control the frequency of an oscillator using another oscillator.

We can expand the patcher by adding some math to make it more usable, as in figure 2.8.

Naive parameters are f_c (Carrier Frequency), f_m (Modulator Frequency), and A_m (modulation Amount).

Naive Implementation

Modulator Frequency (fm)

osc~

Modulation Amount (Am)

*~

Carrier Frequency (fc)

+~

osc~

output~

volume dsp

Figure 2.8.: Naive Implementation with Direct Parametrization.

The output frequencies will be

$$f_c \pm n \cdot f_m \tag{2.7}$$

for n being all integers. So the carrier frequency plus and minus integer multiples of the modulation frequency. We get many (theoretically infinitely) many overtones with different amplitudes this way.

The amplitudes of the different frequencies are determined by Bessel functions which makes it so seemingly random.

Typically, FM is controlled via *Index*, *Ratio*, and fundamental Frequency. The Index, I is given by Modulation Depth and Modulator Frequency.

$$I = \frac{A_m}{f_m} \tag{2.8}$$

A more controllable Implementation will generate the naive parameters from a Ratio, R, the Carrier Frequency and the Index:

$$f_m = \frac{f_c}{R} \tag{2.9}$$

$$A_m = \frac{I}{f_m} \tag{2.10}$$

Figure 2.9.: FM with Index and Ratio

Figure 2.10.: Modulator frequency: 1Hz, Crarrier Frequency 50 Hz, Modulation Amount 10. Please ignore the labeling of the Y axis of the spectrum plot(10^0). It is wrong. The y axis goes from 0 to nyquist=22050Hz.

Figure 2.11.: Modulator frequency: 100Hz, Crarrier Frequency 1000 Hz, Modulation Amount 0.5. Please ignore the labeling of the Y axis of the spectrum plot(10^0). It is wrong. The y axis goes from 0 to nyquist=22050Hz.

2.3. Key Points

- Make sure you understand the differences between AM and FM

- Make sure you recognize what form of modulation is present if you see a patcher or a simplfied formula (like: $sin(a) \cdot sin(b)$ what form of modulation is this?)

- make sure you could make such a simplified formula if you see a patch and vice versa.

- make sure you know what frequencies result from an amplitude modulation

- make sure you know what frequencies result from a frequency modulation

3. Filter

Figure 3.1.: Caravaggio, Narziss. ad Feedback: Echo liebt Narziss etc.

lecture mind map missing

What is a filter? We use filters a lot. We often need to shape the spectrum of a signal. Be it for technical reasons (DC-Offset removal, Anti-aliasing filters, crossovers, interpolation, ...) or for aesthetical ones (EQ-ing, subtractive synthesis)

Please take a moment and seriously ask yourself, knowing what you know about signals, knowing what we can do with a computer: How do we actually do this? What is a digital Filter? This is what this chapter will try to explain.

3.1. Riddle and Getting a Feel for Filtering

Three Questions in order to make you think about signals:

Question 4 *Please look at the image below. Can you draw an approximation of its spectrum? What signals could have been mixed to obtain this result?*

Figure 3.2.: Original Signal

Answer 4 *You should be able to see that there is a sinusoidal component in the signal. This sinusoid has a period of 10 ms, so a frequency of 100 Hz. But also you should see that it is not a clean sine wave, there is some noisy component in there. This is what you should have been able to infer from the image. In fact it is an (attenuated) 100Hz sinusoid mixed with (attenuated) white noise.*

Question 5 *The signal in figure 3.2 was filtered. As a result we obtained the signal in figure 3.3. What kind of filter could have been used?*

Figure 3.3.: Filtered Signal 1

Again, the signal in figure 3.2 (so the original signal again) was filtered. As a result we obtained the signal in figure 3.4 . What kind of filter could have been used?

Question 6

Figure 3.4.: Filtered Signal 2

Woraus könnte das signal in Bild 3.2 bestehen? **Kann jemand das spektrum aufmalen?**

List of Figures

Bibliography

Farnell, A. (2010). *Designing sound.* MIT Press, Cambridge, Mass.